
Report on Solar Potential Prediction using Image Segmentation and Structured Data

Report

Stefan Wezel

stefan.wezel@student.uni-tuebingen.de
4080589
ML4S

Abstract

Solar panels are a cost-effective solution for generating energy in a carbon-free manner. However, not every roof is suitable for installing solar panels. Architecture and location heavily affect the viability of such systems. Predicting this solar potential of a roof is traditionally a labor-intensive process requiring on-site measurements. Automating this process and scaling it up is a difficult challenge. Here, we will introduce a solution proposed by Bacry et al. [1]. We will review their method, examine reported results, and propose possible enhancements.

Introduction

In the European Union (EU) alone, rooftops make up an estimated area of 7935 km² [2]. Much of this area could be used to install solar panels and help feed demand for renewably generated energy. Before installing solar panels, it is useful to estimate how much energy they could produce. This is referred to as a roof's solar potential¹ and is a crucial task. Locally, to determine the viability and economic efficiency of solar panels. Globally, it could also help yield a guess of how much solar energy could contribute to overall energy production capabilities. An estimate of solar potential on a per-country basis of EU members states is illustrated in Figure 1. It highlights the high solar potential in central and southern Europe.

Traditionally, a roof's solar potential is estimated by performing measurements of roof geometry, considering its geographic location, and architecture of surrounding buildings or vegetation [3]. While more recently, geographic information systems (GIS) play an increasingly large role in guiding solar development, much of the process is still labor- and time-consuming. Thus, solar potential estimation on a large scale remains challenging.

Machine learning offers promising capabilities to increase the magnitude at which solar potential estimation can be performed. However, due to limited and complex data, it is not a trivial problem. A solution is proposed by Bacry et al. [1]. They incorporate structured data and existing knowledge as an inductive bias to a method that combines machine learning, algorithmic, and analytical methods. In this work, we will introduce their approach, discuss it, and propose further enhancements to it.

Related work

Freitas et al. [3] present an overview of approaches combining algorithms and GIS modeling to estimate solar potential in densely populated urban environments. They compare different numeric

¹The term technical potential is used interchangeably in this work. In this context, both refer to what is 'technically' possible without economical considerations.

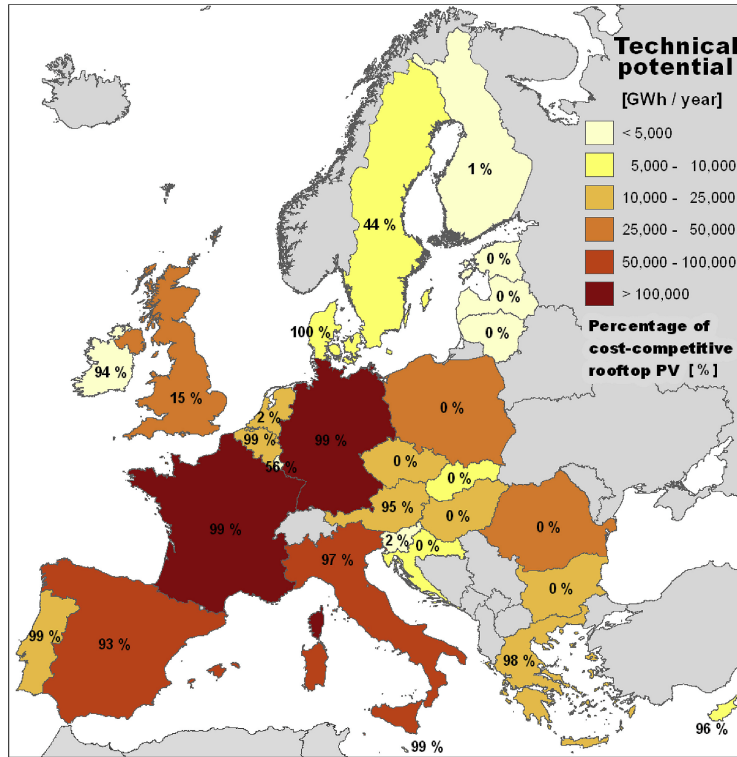


Figure 1: Technical solar potential in GWh/year across EU member states. The percentages indicate the share of cost-effective solutions among all possible solar installations. This accounts for energy prices and the cost of capital. Most countries in central and southern Europe show high technical solar potential exceeding 25 000 GWh/year. A large east-west divide in cost-effectiveness can clearly be observed. France and Germany have high potential and high cost-effectiveness making them primary targets for installing solar panels. The figure is taken from Bódis et al. [2].

solar radiation algorithms and data sources ranging from 2d maps to high-resolution 3d models of urban scenes. In their survey, the authors find that major factors limiting these approaches include poor data quality and the difficulty of validating models.

Bódis et al. [2] use high-resolution satellite data and statistical information to produce an estimate of solar potential across the whole EU. They also include economical calculations in their method to estimate viability of solar installations. However, their method only yields estimates for wider areas and not for individual rooftops.

Early approaches incorporating machine learning are proposed by Ouammi et al. [4] and Sözen et al. [5] who use rudimentary neural network architectures and focus on Moroccan and Turkish territory respectively. Similarly, Assouline et al. [6] focus on Switzerland and use random forests for their predictions.

With project sunroof, the technology company Google has proposed an approach that offers fine-grained solar potential estimation for individual rooftops within the United States and Puerto Rico. From Google Maps data, they find rooftop outlines using a (not further specified) deep learning method. The authors then estimate rooftop geometry and use historical weather data to predict the solar potential [7].

Further private sector endeavors include joint effort between the companies Otovo and In Sun We Trust that offer a product similar to Project Sunroof but only serve France [8]. Other existing products focus on small areas or only produce solar potential estimates on-demand [8, 9].

Lee et al. [10] propose a data-driven method that mostly relies on widely available satellite data. The authors estimate roof topology directly from this imagery using image segmentation architectures. They then use further public data of solar irradiance to estimate solar potential. They validate their method by comparing it to a precise but expensive LIDAR-based approach. Their method can be applied in a wide variety of settings.

Several other approaches leveraging sophisticated deep learning methods are proposed for solar potential estimation [11, 12] or solar irradiance mapping [13, 14, 15].

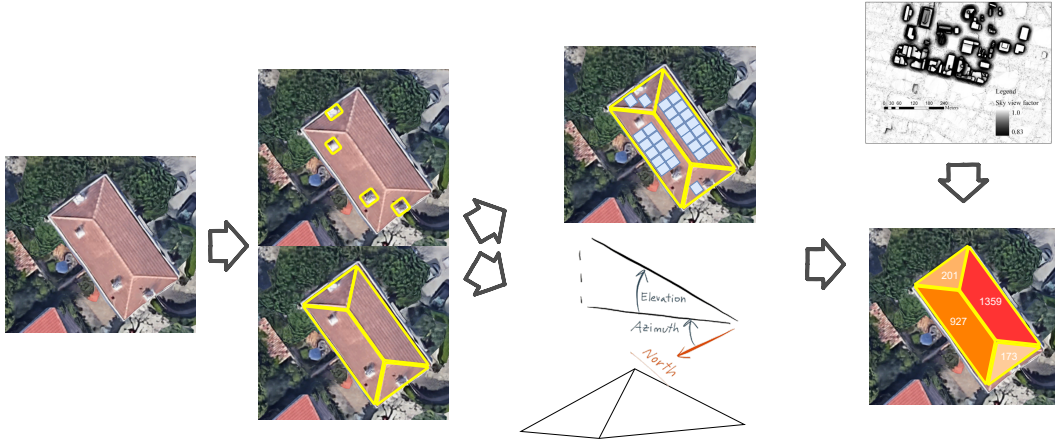


Figure 2: Modular pipeline consisting of segmentation, geometry estimation, panel layout, shading estimation, and eventual solar potential estimation. The figure is partly courtesy of MaxarTechnologies [16] and Zakšek et al. [17].

Method

With the challenges of large-scale solar potential for individual rooftops established and several related proposed methods presented, we will use this section to explain the method proposed by Bacry, de Barros Soares, Andrieux, Hell, Lenhardt, Badosa, Gavaille, and Gaïffas [1] in detail. Its major steps are illustrated in Figure 2.

Segmentation of Aerial Imagery

Given satellite imagery, Bacry et al. [1] first identify suitable roof spaces. This includes finding roof and more precisely, sections of roofs where solar panels can be installed. This means omitting ridges or occupied roof spaces. Therefore, they segment into the categories background, usable segments, and ridges. A second model with the same architecture segments roof objects and predicts a map for classes like chimney, window, and others. They use two different models to achieve this [1].

To extract features from raw images, they use a ResNet backbone [18] with 34 residual blocks.

Inspired by U-Net [19] they store features along the layers during encoding the input. When decoding the features they progressively concatenate these kept features with upsampled ones that match in dimension. This ensures that information from input images is not lost and the model rather has to learn a difference function to the input image as opposed to a whole new image. Satellite imagery is particularly rich in structural information, such as sharp and straight outlines of roofs or roads, making this U-Net style architecture very applicable.

To train the model, Bacry et al. [1] use labels obtained from GIS data and manually annotated images from the French cities of Bordeaux, Brest, Montpellier, and Strasbourg.

Estimating Roof Geometry

Once they obtain the segmentation maps, where roof space suitable for solar panel installation is denoted as a specific class, Bacry et al. [1] predict pitch and azimuth of such areas. This gives information about the geometry of the roof.

The pitch of a roof segment refers to its slope. This plays a crucial role for solar potential since roofs that are either completely flat or vertical might be less beneficial than moderately sloped ones. Bacry et al. [1] regress on the normalized pitch. For this, they use a simple linear regression algorithm [20]. Additionally, Bacry et al. [1] use a random forest algorithm [21] to predict roof inclination where features, such as roof type, material, and shape, or the height of the building are known. The training data is obtained again from 3d models of the five French cities.

The azimuth refers to the orientation of the roof from a bird's view perspective. This is obviously important as i.e. on the northern hemisphere south-facing segments might receive more solar radiation

Table 1: Results for the different subtasks as reported by Bacry et al. [1].

| Task | Method | Metric | Score |
|----------------------------|----------------------|---------------------|-------|
| Roof sections segmentation | ResNet-34-based UNet | Pixel accuracy | 77% |
| Roof objects segmentation | ResNet-34-based UNet | Pixel accuracy | 30% |
| Roof sections orientation | Analytic | Accuracy | 79% |
| Roof sections pitch | Linear regression | Mean absolute error | 3.9° |

than otherwise oriented ones.

Bacry et al. [1] compute the azimuth analytically. For this, they make the simplifying assumption that roof segment orientations can be assigned one of four classes which are the elements of a cyclic 90° rotation group, corresponding to north-, east-, south-, or west-facing segments. This allows them to treat it as a classification problem to which the solution can be obtained in closed form.

Optimal Panel Arrangement

Based on the predicted roof section that is suitable for solar panel installation and its predicted geometry, Bacry et al. [1] compute the maximum number of panels that can be installed. Given the simple rectangular shape of solar panels, they are able to apply a greedy algorithm to the problem. The algorithm’s primary aim is to fit as many panels of a fixed size to a roof segment. It accounts for objects occupying rooftop space, mandatory distances between panels and roof edges, and overlapping.

Shading Estimation

A roofs shading depends on factors such as surrounding vegetation and buildings, its location and geometry, and meteorological factors. To compute the shading, a roof segment is objected to, Bacry et al. [1] use the SkyViewFactor software [17] to obtain data on how much light could reach the panels theoretically. This information is refined further using the R-package shadow [22]. Additionally, Bacry et al. [1] use digital elevation models and projected shadow computations using a QGIS extension [23, 24] to produce a second shading mask per panel. Bacry et al. [1] then combine the two shading masks for a final shading prediction.

Using the shading information, they predict photovoltaic power PV_{DC} . This is converted to PV_{AC} considering losses that might occur in converters or other system components. This value summed up over the span of a year forms PV_{out} .

Solar potential prediction

Finally, Bacry et al. [1] compute solar potential based on estimated number of solar panels and PV_{out} . Additionally, P_{max} is required. This value depends on the solar panel itself and can be treated as a constant. The resulting equation is

$$solar\ potential(kWh/year) = N_{modules} \cdot P_{max} \cdot P_{out}$$

Results & Discussion

With their proposed methods, Bacry et al. [1] are able to estimate solar potential for a great number of roofs. Their method relies heavily on manual data wrangling and steps are performed separately. Training is not done in an end-to-end manner.

While evaluation is possible on hold-outs of the respective datasets, it is difficult to verify whether the datasets’ distribution matches the real world. Results on such hold-out sets are listed in Table 1. As a sanity check, Bacry et al. [1] test whether south-facing roof segments have higher solar potential than north-facing ones and find that this is indeed true. Moreover, they find that east and west-facing segments do not significantly differ in solar potential, which is expected.

Since the proposed method relies so heavily on different data sources it is not trivial to scale to other areas, where a lot of the data used might not be available. Also, architecture or zoning differs from

the French cities the method was trained on. In, for example, Asian cities, flat roofs are much more common which might cause bad segmentation results. Also, roof objects like air conditioners that are not so common in central Europe, might not be identified correctly in that case.

For segmenting the roof sections, Bacry et al. [1] report a pixel accuracy of 77 %. This is a very good value given the challenging task. It should be noted that Bacry et al. [1] do not report mean intersection over union scores (mIoU) which is a more common metric in the image segmentation literature [25]. The pixel accuracy for roof objects is only 30 %. We hypothesize that making the simplifying assumption that there are just two classes for this task, either rooftop object or background would result in a higher pixel accuracy as more capacity in the model’s filters could be dedicated to making the distinction between the two.

Further, we hypothesize that using a unified architecture for segmenting would bring performance benefits. As rooftop segment and rooftop object are mutually exclusive, sharing weights might yield a performance increase and faster convergence. A single segmentation model would then produce segmentation maps with classes rooftop segment, ridge, rooftop object, and background.

As Bacry et al. [1] are likely using cross-entropy as segmentation loss, this approach might also enable a cleaner learning signal as there are no overlapping labels.

Another interesting approach to explore might be to use a rotation equivariant segmentation model. Applying convolutional filters at different orientations could enable further weight sharing. The utilized remote sensing data is particularly suitable for this as the top-down perspective results in the same class of objects occurring in the data, rotated in different ways [26]. A rooftop is a rooftop, no matter how it is rotated.

This might also prove useful for the downstream task of azimuth prediction. If the orientation value is kept throughout the segmentation model, it might not be necessary to discretize it into a classification task but it might be possible to regress on a continuous value.

Bacry et al. [1] state that the segmentation step of the pipeline was particularly challenging [27]. The downstream tasks are more rigid and offer less room for errors. They also report good evaluation results for those.

Conclusion

Bacry et al. [1] propose a method to predict solar potential from satellite imagery and other data such as elevation maps and 3d models. These data are used in separate steps of a pipeline consisting of machine learning models, algorithmic, and analytical steps.

They demonstrate that their method is able to estimate the solar potential of rooftops albeit for selected areas only. We argue that while their method is unlikely to scale beyond the domain it is trained on without further manual labeling, there is potential for more generalization. The modular nature of the pipeline makes it flexible and individual parts might be usable for a wide range of scenarios. Moreover, a more streamlined segmentation model could yield better results for downstream tasks.

References

- [1] Emmanuel Bacry, Daniel de Barros Soares, François Andrieux, Bastien Hell, Julien Lenhardt, Jordi Badosa, Sylvain Gavaille, and Stéphane Gaïffas. Predicting the solar potential of rooftops using image segmentation and structured data. In *NIPS Proceedings*, 2021.
- [2] Katalin Bódis, Ioannis Kougias, Arnulf Jäger-Waldau, Nigel Taylor, and Sándor Szabó. A high-resolution geospatial assessment of the rooftop solar photovoltaic potential in the european union. *Renewable and Sustainable Energy Reviews*, 114:109309, 2019.
- [3] Sara Freitas, Cristina Catita, Paula Redweik, and Miguel Centeno Brito. Modelling solar potential in the urban environment: State-of-the-art review. *Renewable and Sustainable Energy Reviews*, 41:915–931, 2015.
- [4] Ahmed Ouammi, Driss Zejli, Hanane Dagdougui, and Rachid Benchrif. Artificial neural network analysis of moroccan solar potential. *Renewable and Sustainable Energy Reviews*, 16(7):4876–4889, 2012.
- [5] Adnan Sözen, Erol Arcaklioğlu, Mehmet Özalp, and E Galip Kanit. Use of artificial neural networks for mapping of solar potential in turkey. *Applied Energy*, 77(3):273–286, 2004.
- [6] Dan Assouline, Nahid Mohajeri, and Jean-Louis Scartezzini. Large-scale rooftop solar photovoltaic technical potential estimation using random forests. *Applied energy*, 217:189–211, 2018.
- [7] Project sunroof data explorer: a description of methodology and inputs technical report, 2017. URL <https://sunroof.withgoogle.com/>.
- [8] Simulateur panneau solaire, 2021. URL <https://simulateur.insunwetrust.solar/>.
- [9] Rhino solaire, 2021. URL <https://rhinoterrain.com/fr/rhinosolar.html>.
- [10] Stephen Lee, Srinivasan Iyengar, Menghong Feng, Prashant Shenoy, and Subhransu Maji. Deeproof: A data-driven approach for solar potential estimation using rooftop imagery. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 2105–2113, 2019.
- [11] Zhaojian Huang, Thushini Mendis, and Shen Xu. Urban solar utilization potential mapping via deep learning technology: A case study of wuhan, china. *Applied Energy*, 250:283–291, 2019.
- [12] Teng Zhong, Zhixin Zhang, Min Chen, Kai Zhang, Zixuan Zhou, Rui Zhu, Yijie Wang, Guonian Lü, and Jinyue Yan. A city-scale estimation of rooftop solar photovoltaic potential based on deep learning. *Applied Energy*, 298:117132, 2021.
- [13] Pratima Kumari and Durga Toshniwal. Deep learning models for solar irradiance forecasting: A comprehensive review. *Journal of Cleaner Production*, 318:128566, 2021.
- [14] Deeksha Chandola, Harsh Gupta, Vinay Anand Tikkiwal, and Manoj Kumar Bohra. Multi-step ahead forecasting of global solar radiation for arid zones using deep learning. *Procedia Computer Science*, 167:626–635, 2020.
- [15] Olusola Bamisile, Ariyo Oluwasanmi, Sandra Obiora, Emmanuel Osei-Mensah, Gaylord Asoronye, and Qi Huang. Application of deep learning for solar irradiance and solar photovoltaic multi-parameter forecast. *Energy Sources, Part A: Recovery, Utilization, and Environmental Effects*, pages 1–21, 2020.
- [16] Copernicus MaxarTechnologies. Google imagery, geobasis-de bkg, geocontent, landsat / copernicus maxar technologies, 2022. URL <https://www.google.com/maps/>.
- [17] Klemen Zakšek, Kristof Oštir, and Žiga Kokalj. Sky-view factor as a relief visualization technique. *Remote sensing*, 3(2):398–415, 2011.
- [18] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- [19] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical image computing and computer-assisted intervention*, pages 234–241. Springer, 2015.
- [20] Jurgen Gross and Jürgen Groß. *Linear regression*, volume 175. Springer Science & Business Media, 2003.
- [21] Mariana Belgiu and Lucian Drăguț. Random forest in remote sensing: A review of applications and future directions. *ISPRS journal of photogrammetry and remote sensing*, 114:24–31, 2016.
- [22] Michael Dorman, Evyatar Erell, Adi Vulkan, and Itai Kloog. shadow: R package for geometric shadow calculations in an urban environment. *R J.*, 11(1):287, 2019.
- [23] Qgis, 2020. URL <https://qgis.org>.

- [24] qgis-shadows, 2019. URL <https://landscapearchaeology.org/2019/qgis-shadows/>.
- [25] Shervin Minaee, Yuri Y Boykov, Fatih Porikli, Antonio J Plaza, Nasser Kehtarnavaz, and Demetri Terzopoulos. Image segmentation using deep learning: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2021.
- [26] Diego Marcos, Michele Volpi, Benjamin Kellenberger, and Devis Tuia. Land cover mapping at very high resolution with rotation equivariant cnns: Towards small yet accurate models. *ISPRS journal of photogrammetry and remote sensing*, 145:96–107, 2018.
- [27] François Andrieux. Predicting the solar potential of rooftops using image segmentation and structured data, Oct 2021. URL <https://medium.com/nam-r/predicting-the-solar-potential-of-rooftops-using-image-segmentation-and-structur>